


Severe tier



TWO different exclusion marks: the '///' band = excluded for ALL models (agent drop, not poolable); the 'xxx' N/A cell = THIS model could not run it. CoSDH agent_drop and missing_modality are NOT APPLICABLE (6 cells): the released inference path CRASHES when a collaborator is removed, rather than degrading -- which is itself the finding. Shown as N/A, never as 0.0, excluded from every ranking. DATASET CAVEAT: CoSDH runs on OPV2V test (2,170 frames); the other three on V2XSet test (2,834). rel_drop normalises away different clean baselines but not a different dataset -- cross-model comparisons involving CoSDH are NOT like-for-like. Ungraded for different reasons: Where2comm = PROVENANCE (third-party retrain, no official OPV2V/V2XSet release); CoSDH = RETRAIN (author release, but README states the checkpoints were retrained and 'differ slightly' from the paper; measured clean 0.9660/0.9289 vs published 0.9683/0.9299). Agent drop is a hatched band rather than silently dropped: stratified by scene agent count (Fig. 4). LiDAR fog is a normal ranked row: monotone on all four models, no exact zeros (CoBEVT -0.875, V2X-ViT -0.889, Where2comm -0.988, CoSDH -0.743 on OPV2V). Where2comm at severe is AP 0.0067 — near the floor and to be read as near-saturated in isolation, though it ranks distinctly. LiDAR fog PROVENANCE: the 9 V2XSet fog cells were regenerated by the CORRECTED injector (lidar_fog: bdc5c3f7) after the intensity-binarisation fix; CoSDH's 3 fog cells predate driver stamping (fog code byte-identical per adapter digest), making fog the ONE injector the census flags as version-mixed -- do not pool CoSDH fog with the other three. Every non-fog cell predates stamping uniformly.