




CoSDH agent_drop and missing_modality are NOT APPLICABLE (6 cells): the released inference path CRASHES when a collaborator is removed, rather than degrading -- which is itself the finding. Shown as N/A, never as 0.0, excluded from every ranking. DATASET CAVEAT: CoSDH runs on OPV2V test (2,170 frames); the other three on V2XSet test (2,834). rel_drop normalises away different clean baselines but not a different dataset -- cross-model comparisons involving CoSDH are NOT like-for-like. CoSDH therefore contributes to fog, snow, points_reduce, pose and latency only, and the ordering mean on the two lost-collaborator rows runs over the three V2XSet models. GT-DENOMINATOR CAVEAT (agent_drop only): ground truth is the UNION over CAVs, so dropping an agent removes the boxes only it observed. Measured on 200 frames of V2XSet: GT shrinks -2.8% / -5.8% / -8.5% at p=0.25/0.50/0.75, and with EVERY collaborator dropped -12.1%. Recall = TP/GT, so a smaller denominator inflates recall and agent_drop's degradation is biased toward looking Milder than it is. missing_modality is NOT affected -- it empties a cloud but KEEPS the agent, so its denominator is stable. Consequently agent_drop's POSITION RELATIVE TO MISSING LIDAR IS NOT RELIABLE and could plausibly reverse: the gap between them is 0.06/0.03/0.04 while the denominator shift is 8.5% at severe. The BAND-level claim -- lost-collaborator faults are milder than corrupted-data ones, which sit at -0.3 to -0.9 -- is unaffected and is what this figure asserts. *agent_drop: n=2 stratum only (Fig. 4). NOTE V2X-ViT pose error (-0.16) sits BELOW its own latency (-0.19) and near the lost-collaborator band, vs -0.46 (CoBEVT) and -0.40 (Where2comm). V2X-ViT is the only model shipping loc_err=true. Consistent with noise-trained pose defence, but an ASSOCIATION, not a proven cause -- the three models differ on architecture, dataset pairing and compression too. LiDAR fog now ranks first legitimately (monotone, no exact zeros) rather than as the floor tie it was before the fix. LiDAR fog PROVENANCE: the 9 V2XSet fog cells were regenerated by the CORRECTED injector (lidar_fog: bdc5c3f7) after the intensity-binarisation fix; CoSDH's 3 fog cells predate driver stamping (fog code byte-identical per adapter digest), making fog the ONE injector the census flags as version-mixed -- do not pool CoSDH fog with the other three. Every non-fog cell predates stamping uniformly.