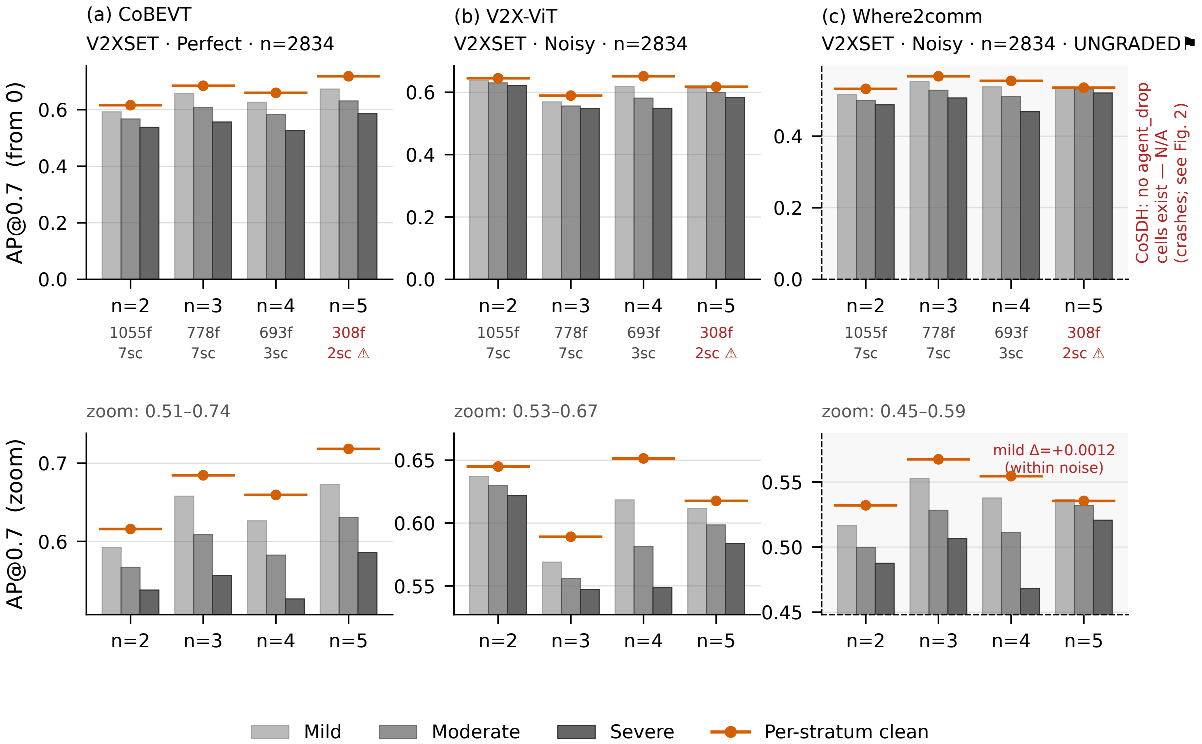

CoSDH IS ABSENT FROM THIS FIGURE AND CANNOT APPEAR: it has no agent_drop cells at all -- the released inference path crashes when a collaborator is removed rather than degrading (manifest not_applicable; Fig. 2 N/A cells). The absence is the datum. GT-DENOMINATOR CAVEAT (agent_drop only): ground truth is the UNION over CAVs, so dropping an agent removes the boxes only it observed. Measured on 200 frames of V2XSet: GT shrinks -2.8% / -5.8% / -8.5% at p=0.25/0.50/0.75, and with EVERY collaborator dropped -12.1%. Recall = TP/GT, so a smaller denominator inflates recall and agent_drop's degradation is biased toward looking MILDER than it is. missing_modality is NOT affected -- it empties a cloud but KEEPS the agent, so its denominator is stable. STRATUM DEPENDENCE: the strata are NOT equally biased. With every collaborator dropped, GT falls -11.2% at n=2 but -22.7% at n=3 -- more collaborators means more of GT is uniquely theirs -- so the bias is LARGEST exactly where collaboration should look most valuable, and it runs in the direction that FLATTERS collaboration. n=4 and n=5 magnitudes are UNMEASURED: those strata did not occur in the 200 frames sampled. Top row: bars from zero (honest baseline). Bottom row: the SAME data zoomed to the working range — the axis is not truncated in the top row, which would exaggerate differences. x-axis is scene agent count (post max_cav truncation); f = ego frames, sc = scenarios; ⚠ thin stratum, unreliable. Per-stratum clean references differ by up to 10.2 AP (CoBEVT n=2 0.616 -> n=5 0.718), which is why a pooled agent-drop number would be meaningless. At n=2, dropping a collaborator is ablation to single-agent perception, not graded degradation.